

VIRTUALIZATION: THE CPU

Andrea Arpaci-Dusseau
CS 537, Fall 2019

ADMINISTRIVIA

- Project 1 is out! Due Monday, 9/16 before midnight
 - Solo, but later ones will involve project partners
 - Handin directories and test cases available
- Discussion sections on Wednesday
 - Can attend others if room (last two have most space)
- Sign up for Piazza
- Lecture recordings
- Still on waitlist? Sign attendance sheet at end of lecture
- Lecture ends at 12:15pm
- Microphone: Let me know if you can't hear

AGENDA / OUTCOMES

Abstraction for CPU

What is a Process? What is its lifecycle?

Mechanism

How does a process interact with the OS?

How does the OS switch between processes?

REVIEW

What is an Operating System?

- Software that converts hardware into a useful form for applications

What abstraction does the OS provide for the CPU?

- Process or thread

For memory?

- Virtual address space

What are some advantages of OS providing resource management?

- Protect applications from one another
- Provide fair and efficient access to resources (cost, time, energy)

VIRTUALIZING THE CPU

High-level Goal:

Give each “process” impression it alone is actively using CPU

Resources can be shared in **time** and/or **space**

Assume single uniprocessor

- Time-sharing (multi-processors: advanced issue with space-sharing)

Memory?

- Space-sharing (later)

Disk?

- Space-sharing (later)

ABSTRACTION: PROCESS

PROGRAM VS PROCESS

```
#include <stdio.h>
#include <stdlib.h>
#include "common.h"
```

```
int main(int argc, char *argv[]) {
    char *str = argv[1];
```

```
    while (1) {
        printf("%s\n", str);
        Spin(1);
```

```
    }
    return 0;
```

```
}
```

Static



Program

Running



Process

WHAT IS A PROCESS?

Stream of executing instructions and their “context” in address space

Instruction
Pointer →


Stack pointer

```
pushq    %rbp
movq     %rsp, %rbp
subq     $32, %rsp
movl     $0, -4(%rbp)
movl     %edi, -8(%rbp)
movq     %rsi, -16(%rbp)
cmpl     $2, -8(%rbp)
je       LBB0_2
```

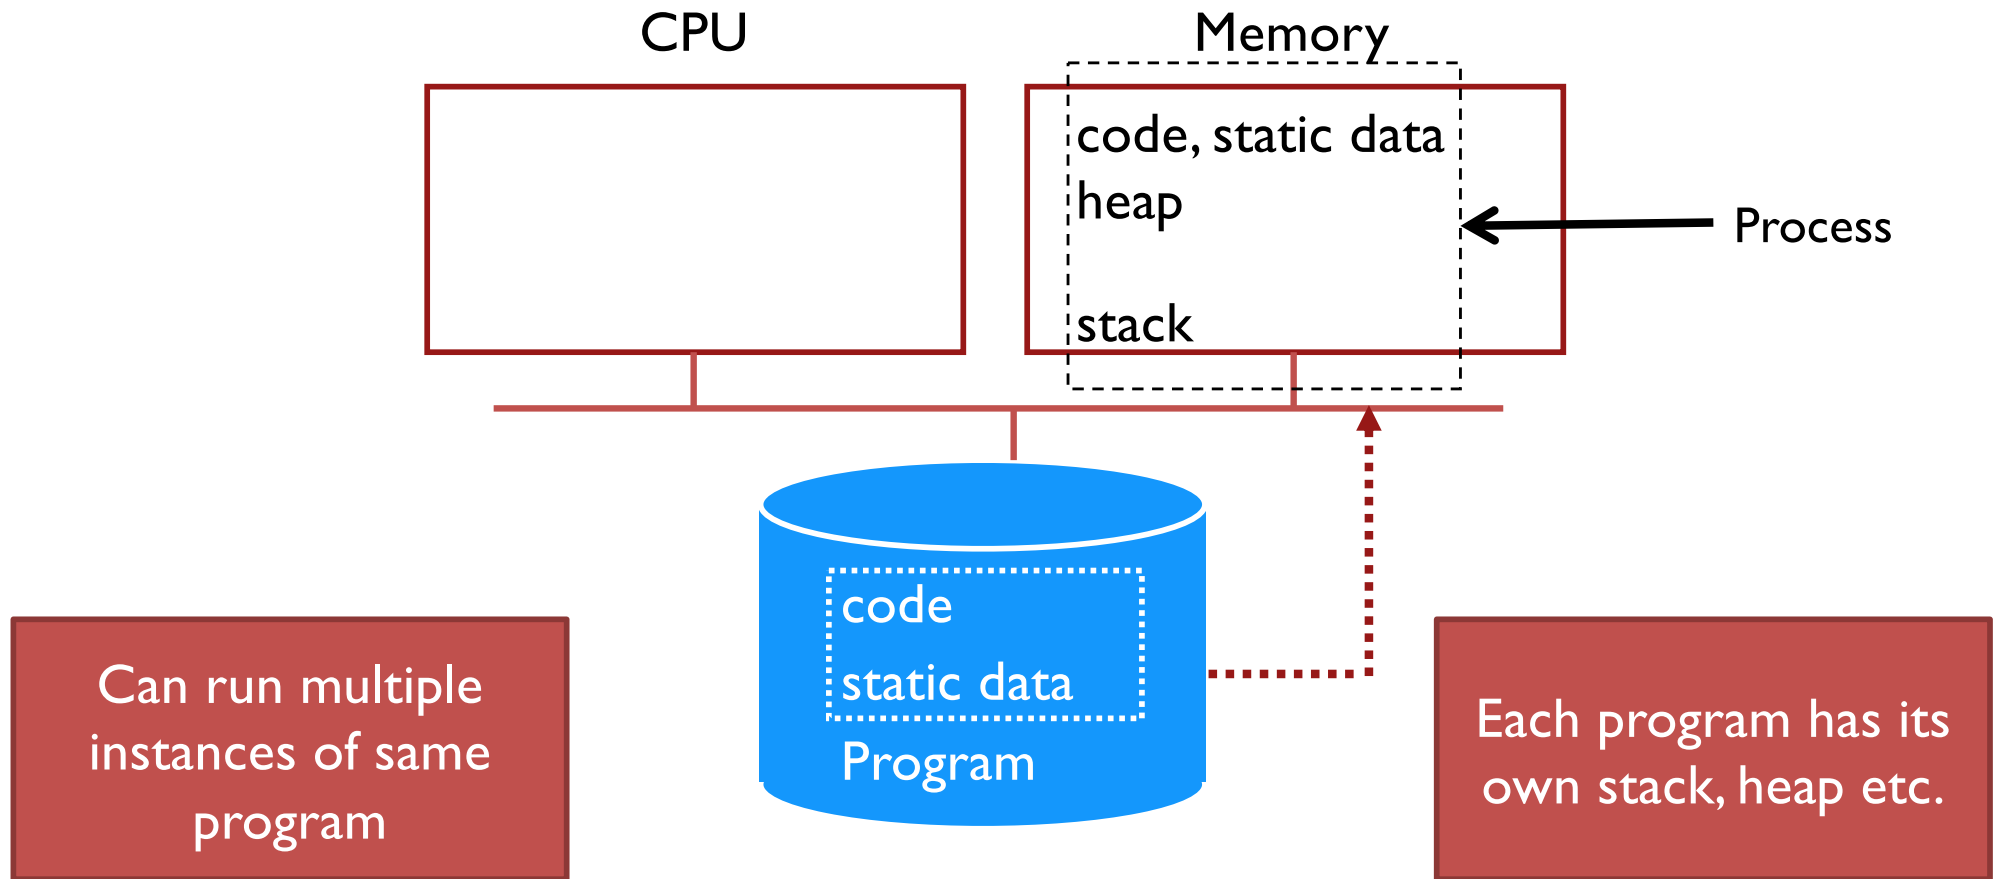
Registers
Memory addrs

File descriptors

PROCESS CREATION



PROCESS CREATION



PROCESS VS THREAD DEMO

- Two **processes** examining same memory address see **different** values (i.e., different contents)
 - Different isolated address spaces
- Two **threads** examining memory address see **same** value (i.e., same contents)
 - Share same address space

PROCESS VS THREAD

Threads: “Lightweight process”


Execution streams that share an address space

Can directly read / write same memory

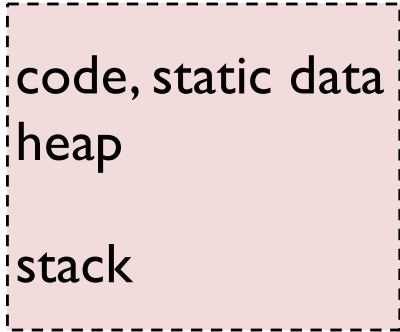
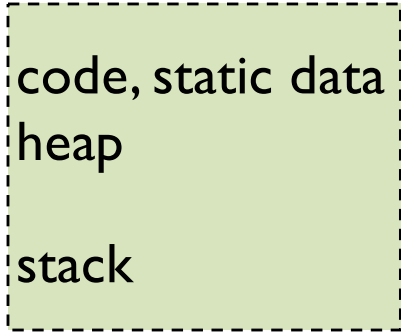
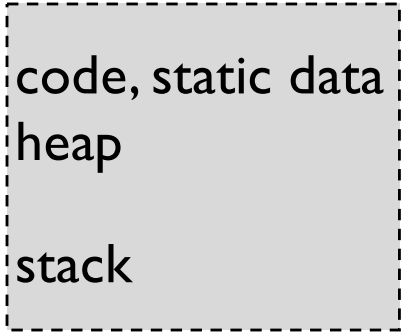
Can have multiple threads within a single process

WHY DO WE NEED PROCESSES ?


SHARING CPU



TIME SHARING




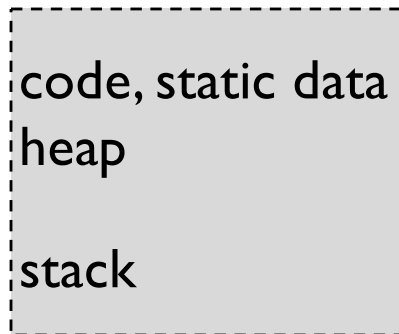


Context is loaded into CPU



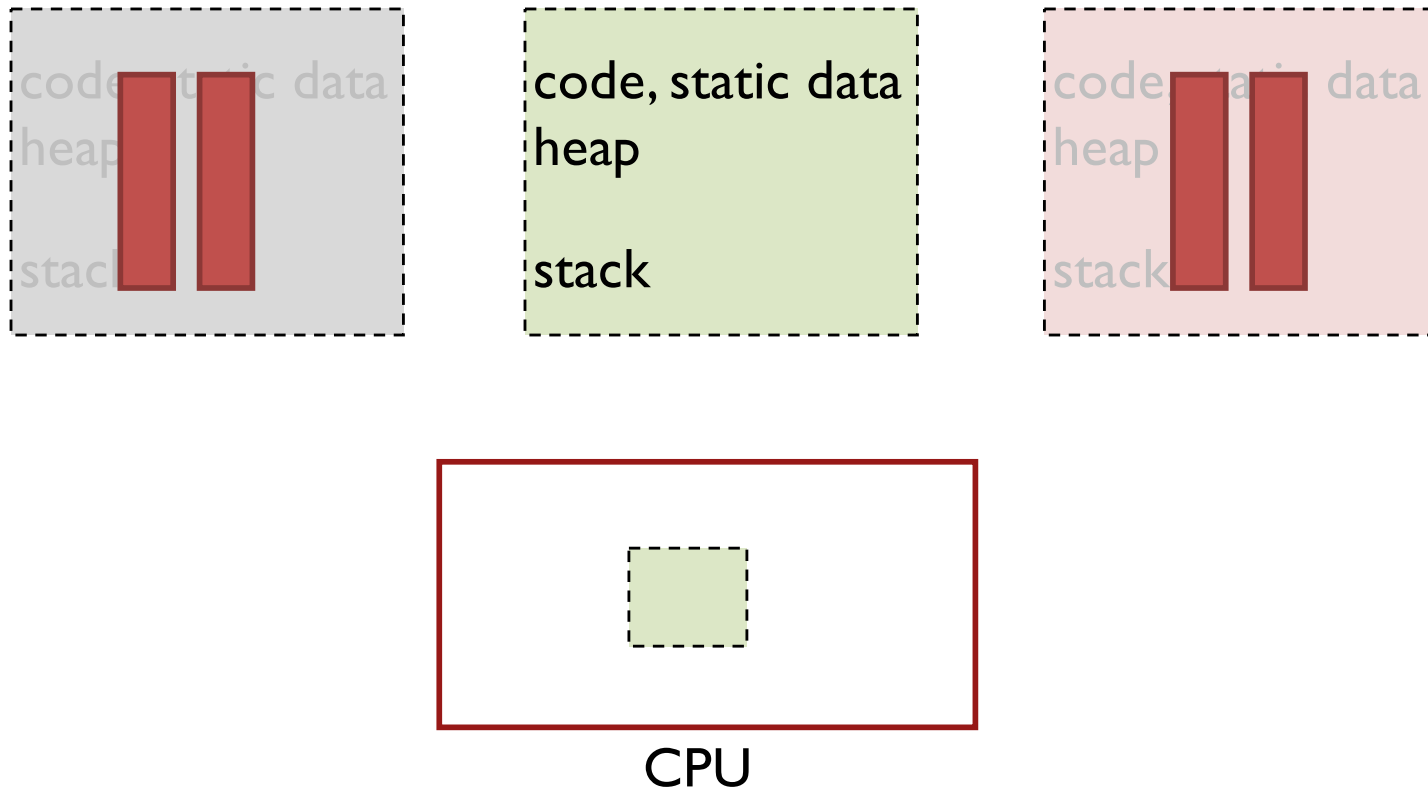
CPU

TIME SHARING



CPU

TIME SHARING



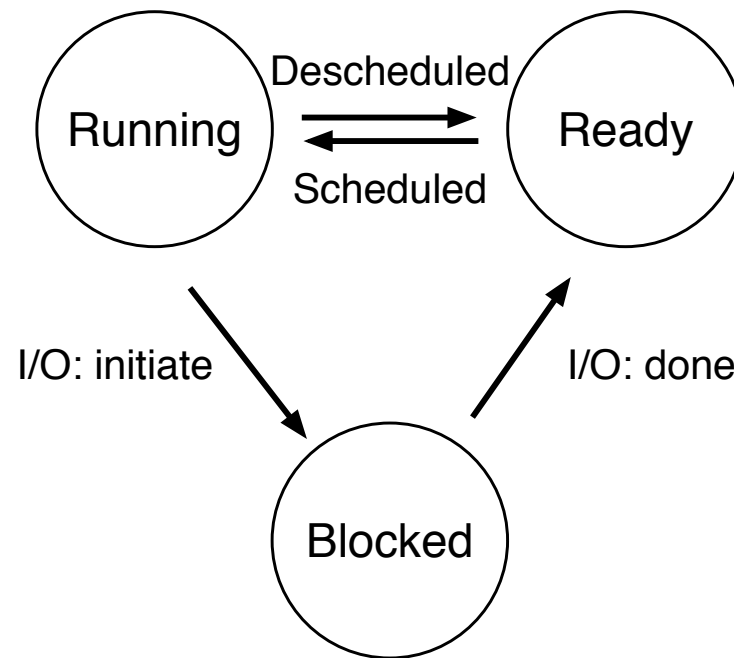
WHAT TO DO WITH PROCESSES THAT ARE NOT RUNNING ?

OS Scheduler

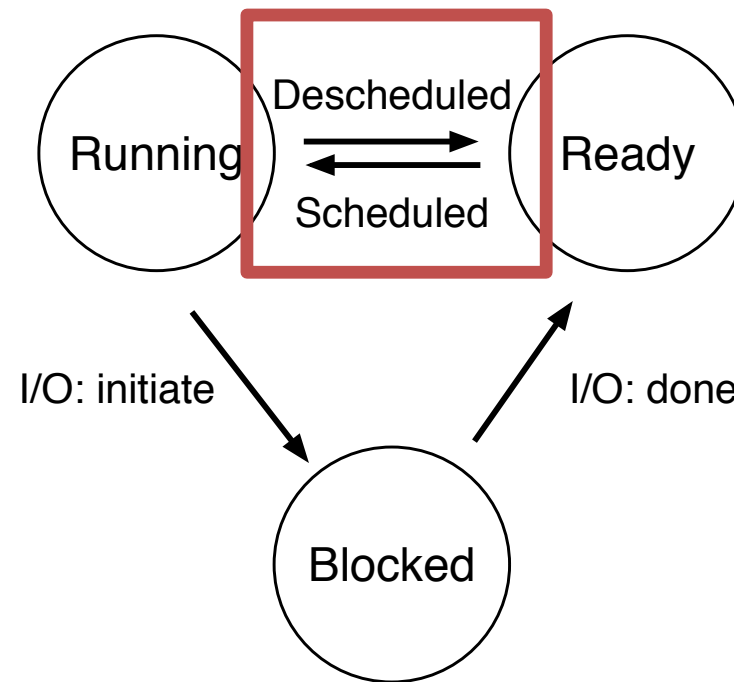
Save **context** when process is paused

Restore context on resumption

STATE TRANSITIONS



STATE TRANSITIONS



ASIDE: OSTEP HOMEWORKS!

- Optional homeworks corresponding to each chapter in book
- Little simulators to help you understand
- Can generate problems and solutions!

<http://pages.cs.wisc.edu/~remzi/OSTEP/Homework/homework.html>

PROCESS HW

Run `./process_run.py -l 2:100,2:0`

QUIZ

\geq ./process-run.py -l 3:50,3:40

Process 0

io

io

cpu

Process 1

cpu

io

io

CPU TIME SHARING

Mechanism goals: Be able to run processes

Efficiency: Time sharing should not add overhead

Control: OS should be able to intervene when required

Policy goals: Pick the “best” process to schedule

Reschedule process for fairness? efficiency ?

Separate mechanism from policy for clean OS design

How to have efficient mechanism??

EFFICIENT EXECUTION

Simple answer !?: **Direct Execution**

Allow user process to run directly

Create process and transfer control to main()

Challenges

- 1) What if the process wants to do something restricted ? Access disk ?
- 2) What if the process runs forever ? Buggy ? Malicious ?

Solution: **Limited Direct Execution (LDE)**

CHALLENGE 1: RESTRICTED OPS

How can we ensure user process can't harm others?

Solution: privilege levels supported by hardware (bit of status)

User processes run in user mode (restricted mode)

OS runs in kernel mode (not restricted)

How can process access devices?

System calls (function call implemented by OS)

SYSTEM CALL

SYSTEM CALL



P wants to call read()

SYSTEM CALL



P can only see its own memory because of **user mode**
(other areas, including kernel, are hidden)

SYSTEM CALL




P wants to call `read()` but no way to call it directly

SYSTEM CALL



```
movl $6, %eax;    int $64
```

SYSTEM CALL




`movl $6, %eax;`

`int $64`

Trap table
index


SYSTEM CALL



Syscall table
index

`movl $6, %eax; int $64`

SYSTEM CALL




Syscall table
index

`movl $6, %eax;`

`int $64`

Trap table
index


SYSTEM CALL



```
movl $6, %eax;    int $64
```

Follow entries to correct system call code

SYSTEM CALL



Kernel can access user memory to fill in user buffer
return-from-trap at end to return to Process P

SYSCALL SUMMARY

Separate user-mode from kernel mode for security

Syscall: call kernel mode functions

Transfer from user-mode to kernel-mode (trap)

Return from kernel-mode to user-mode (return-from-trap)

5 MINUTE BREAK!

Talk with at least two neighbors

- What has been your favorite course in CS?
What did you like most about it?
- Favorite course outside of CS? Why?

REPEAT: EFFICIENT EXECUTION

Simple answer !?: **Direct Execution**

Allow user process to run directly

Create process and transfer control to main()

Challenges

1) What if the process wants to do something restricted ? Access disk ?

2) What if the process runs forever ? Buggy ? Malicious ?

Solution: **Limited Direct Execution (LDE)**

CHALLENGE 2: HOW TO TAKE CPU AWAY

Policy

- To decide which process to schedule when

- Decision-maker to optimize some workload performance metric

Mechanism

- To switch between processes

- Low-level code that implements the decision


Separation of policy and mechanism: Recurring theme in OS

DISPATCH MECHANISM

OS runs **dispatch loop**

```
while (1) {  
    run process A for some time-slice  
    stop process A and save its context  
    load context of another process B  
}
```

Context-switch




Question 1: How does dispatcher gain control?

Question 2: What must be saved and restored?

HOW DOES DISPATCHER GET CONTROL?

Option 1: **Cooperative Multi-tasking**: Trust process to relinquish CPU through traps

- Examples: System call, page fault (access page not in main memory), or error (illegal instruction or divide by zero)
- Provide special `yield()` system call



PROBLEMS WITH COOPERATIVE ?

Disadvantages: Processes can **misbehave**

By avoiding all traps and performing no I/O, can take over entire machine

Only solution: Reboot!

Not performed in modern operating systems

TIMER-BASED INTERRUPTS

Option 2: Timer-based Multi-tasking (True multi-tasking)

Guarantee OS can obtain control periodically

Enter OS by enabling periodic alarm clock

Hardware generates timer interrupt (CPU or separate chip) Example: Every 10ms

User must not be able to mask timer interrupt

Operating System

Hardware

Program
Process A

Operating System

Hardware

Program
Process A

timer interrupt
save regs(A) to k-stack(A)
move to kernel mode
jump to trap handler for timer

Operating System

Hardware

Program
Process A

timer interrupt
save regs(A) to k-stack(A)
move to kernel mode
jump to trap handler for timer

Handle the trap for timer
Call switch() routine
save regs(A) to proc-struct(A)
restore regs(B) from proc-struct(B)
switch to k-stack(B)
return-from-trap (into B)

Operating System

Hardware

Program Process A

Handle the trap for timer
Call switch() routine
save regs(A) to proc-struct(A)
restore regs(B) from proc-struct(B)
switch to k-stack(B)
return-from-trap (into B)

timer interrupt
save regs(A) to k-stack(A)
move to kernel mode
jump to trap handler for timer

restore regs(B) from k-stack(B)
move to user mode
jump to B's IP

Operating System

Hardware

Program Process A

Handle the trap for timer
Call switch() routine
save regs(A) to proc-struct(A)
restore regs(B) from proc-struct(B)
switch to k-stack(B)
return-from-trap (into B)

timer interrupt
save regs(A) to k-stack(A)
move to kernel mode
jump to trap handler for timer

restore regs(B) from k-stack(B)
move to user mode
jump to B's IP

Process B

SUMMARY


Process: Abstraction to virtualize CPU

Use time-sharing in OS to switch between processes

Key aspects

- Use system calls to run access devices etc. from user mode

- Context-switch using interrupts for multi-tasking



POLICY ?
NEXT CLASS!

ADMINISTRIVIA

- Lecture ends at 12:15pm
- Project 1 is out! Due Monday, 9/16 before midnight
 - Handin directories and test cases available
- Discussion sections on Wednesday
 - Can attend others if room (last two have most space)
- Sign up for Piazza
- Still on waitlist? Sign attendance sheet at end of lecture