

# CONCURRENCY: INTRODUCTION

Andrea Arpaci-Dusseau  
CS 537, Fall 2019

# ADMINISTRIVIA

Project 3 Due Tonight

- Turn in Makefile and all src code (.c and .h) to make mysh

Plan Project 4 to be available Tuesday

Office hours instead of Lab Hours from Friday to Monday?

Midterm I on Thursday 7:30pm - 9:30pm

- Covers material including today's lecture
- Read textbook if you haven't
- Complete Canvas homeworks
- Take sample exams and check what you don't understand
- Look over lecture notes; can you answer all questions?
- Post questions for Tuesday lecture review

# AGENDA / LEARNING OUTCOMES

Virtual memory: Summary

Concurrency

What is the motivation for concurrent execution?

What are some of the challenges?

**RECAP**

# SWAPPING INTUITION

Idea: OS keeps unreferenced pages on disk

- Slower, cheaper backing store than memory

Process can run when not all pages are loaded into main memory

OS and hardware cooperate to make large disk seem like memory

- Same behavior as if all of address space in main memory

Requirements:

- OS must have **mechanism** to identify location of each page in address space → in memory or on disk
- OS must have **policy** for determining which pages live in memory and which on disk

# VIRTUAL MEMORY MECHANISMS

First, hardware checks TLB for virtual address

- if **TLB hit**, address translation is done; page in physical memory

Else **TLB miss...**

- Hardware or OS walk page tables
- If PTE designates page is **present**, then page in physical memory (i.e., present bit is cleared)

Else **Page fault...**

- Trap into OS (not handled by hardware)
- OS selects victim page in memory to replace (look at **use** bits for clock)
  - Write victim page out to disk if modified (add **dirty** bit to PTE)
- OS reads referenced page from disk into memory (faulting process is **blocked**)
- Page table is updated, **present** bit is set
- Process continues execution (process moved to **ready** state)



# CHAT: PERFORMANCE IMPACT?

Assume workload of multiple processes

- Processes each alternate between CPU and I/O
- Each access some amount of memory

What happens to system performance (throughput = jobs/sec) as we increase the number of processes?

- If the sum of the working sets  $>$  physical memory?



# BONUS: WHAT IF NO HARDWARE SUPPORT?

What can the OS do if hardware does not have use bit (or dirty bit) (and hardware-filled TLB)?

- Can the OS “emulate” these bits?

Leading question:

- How can the OS get control (i.e., generate a trap) every time use bit should be set? (i.e., when a page is accessed?)



# SUMMARY: VIRTUAL MEMORY

Abstraction: Virtual address space with code, heap, stack

Address translation

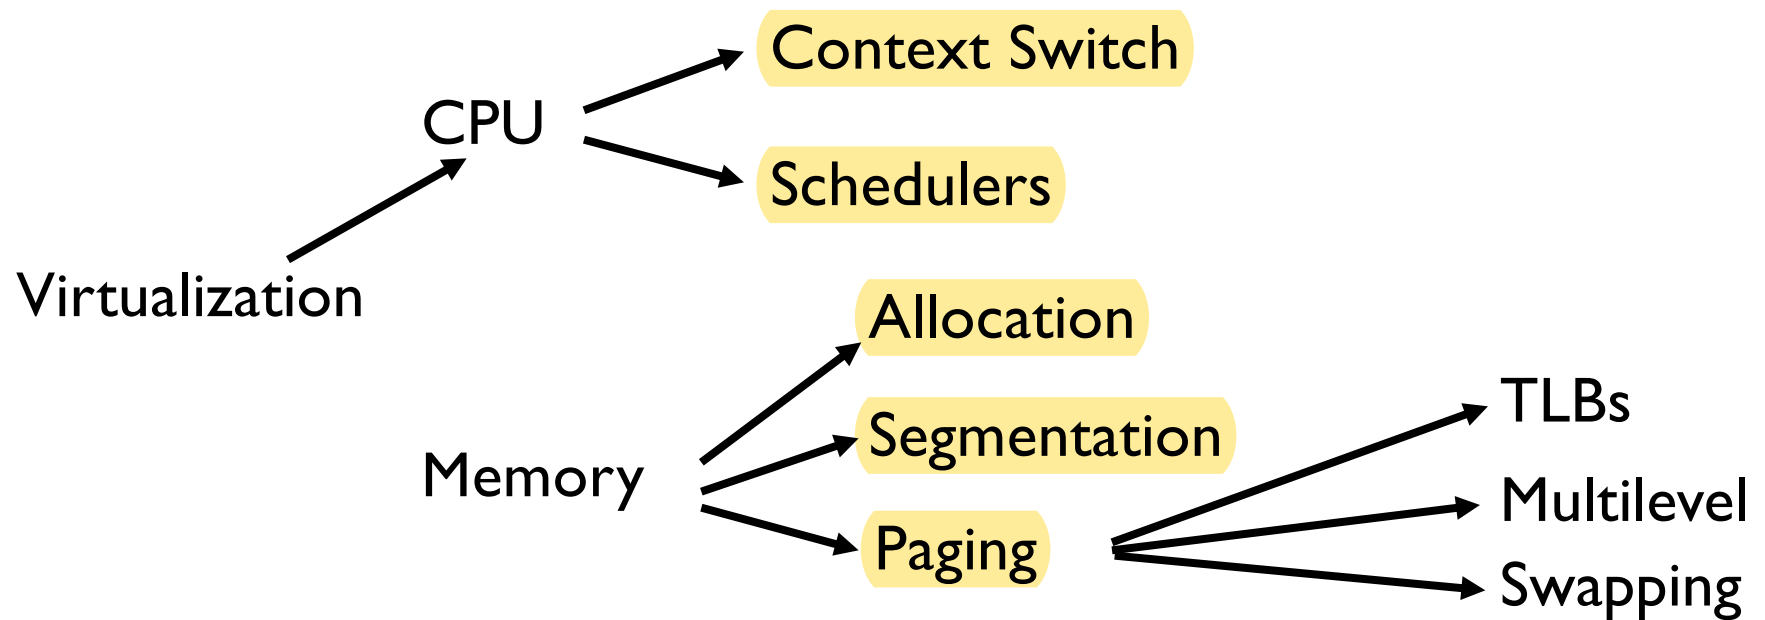
- Contiguous memory: base, bounds, segmentation
- Using fixed sizes pages with page tables

Challenges with paging

- Extra memory references: avoid with TLB
- Page table size: avoid with multi-level paging


Larger address spaces: Swapping mechanisms, policies (LRU, Clock)

# REVIEW: EASY PIECE 1



**CONCURRENCY**

# MOTIVATION FOR CONCURRENCY



# MOTIVATION

CPU Trend: Same speed, but multiple cores

**Option 0:** Run many different applications on one machine  
– OS scheduler ensures cores are used efficiently

Our Goal: Write applications that fully utilize many cores

# OPTION 1: USE MULTIPLE PROCESSES

**Option 1:** Build apps from many communicating **processes**

- Example: Chrome (process per tab)
- Communicate via `pipe()` or other message passing primitives

Pros?

- Don't need new abstractions; good for security

Cons?

- Cumbersome programming
- High communication overheads
- Expensive context switching (why expensive?)

## OPTION 2: DIVIDE PROCESS INTO THREADS

New abstraction: **thread**

Threads are like processes, except:

**multiple threads of same process share an address space**

Divide large task across several **cooperative** threads

Communicate through shared address space

# COMMON PROGRAMMING MODELS

Multi-threaded programs tend to be structured as:

- **Producer/consumer**

Multiple producer threads create data (or work) that is handled by one of the multiple consumer threads

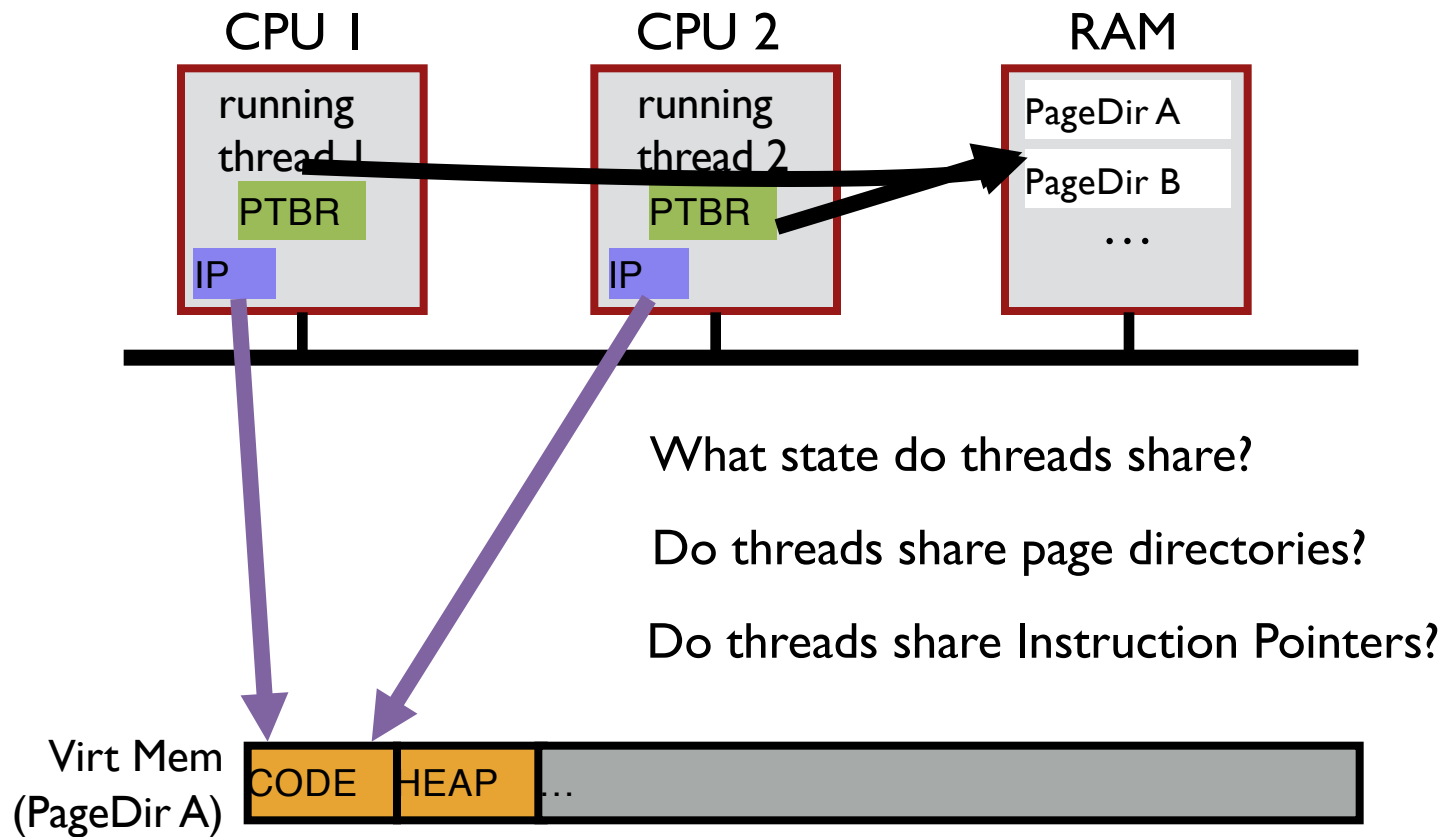
- **Pipeline**

Task is divided into series of subtasks, each of which is handled in series by a different thread

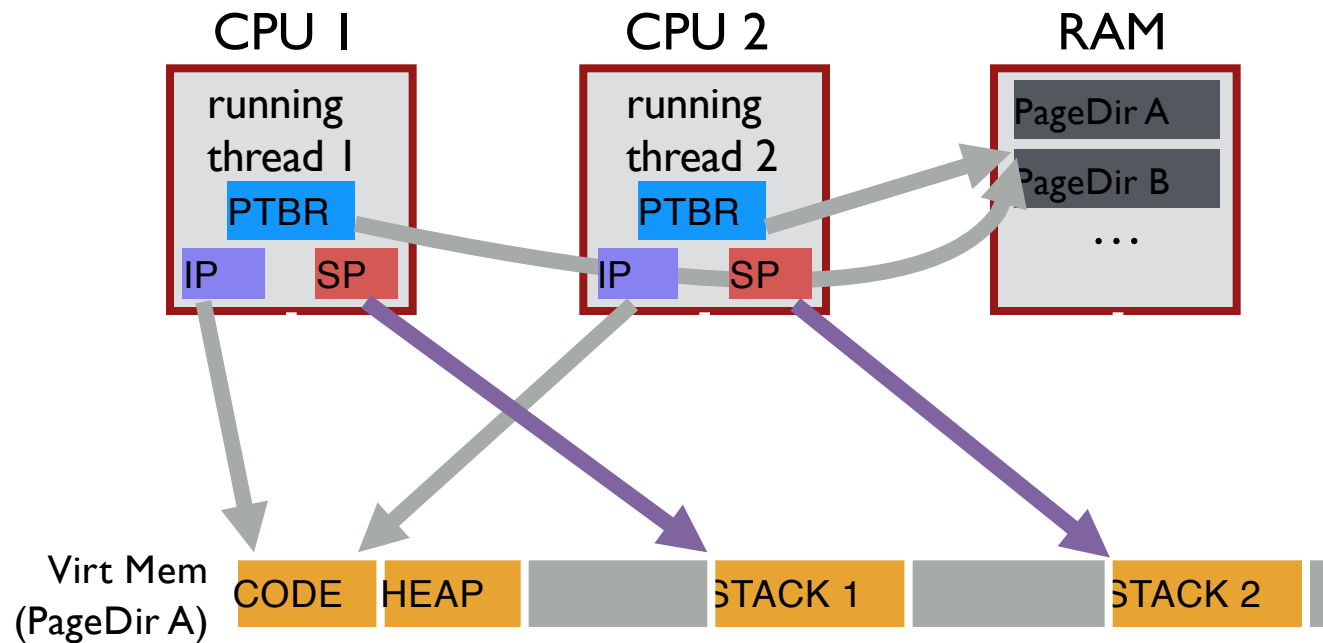
- **Defer work with background thread**

One thread performs non-critical work in the background (e.g., when CPU idle)





Share code, but each thread may be executing different code at the same time  
→ Different Instruction Pointers



Do threads share stack pointer?

Threads executing different functions need different stacks  
(But, stacks are in same address space, so trusted to be cooperative)

# THREAD VS. PROCESS

Multiple threads within a single process share:

- Process ID (PID)
- Address space: Code (instructions), Most data (heap)
- Open file descriptors
- Current working directory (environment variables)
- User and group id (permissions, limits)

Each thread has its own

- Thread ID (TID)
- Set of registers, including Program counter and Stack pointer
- Stack for local variables and return addresses (in same address space)

# THREAD API

Variety of thread systems exist

- POSIX Pthreads

Common thread operations

- Create
- Exit
- Join (instead of wait() for processes)

# OS SUPPORT: APPROACH 1

## User-level threads: Many-to-one thread mapping

- Implemented by user-level runtime libraries
- Create, schedule, synchronize threads at user-level
- OS is not aware of user-level threads
- OS thinks each process contains only a single thread of control

### Advantages

- Does not require OS support; Portable
- Can tune scheduling policy to meet application demands (simplify)
- Lower overhead thread operations since no system call

### Disadvantages?

- Cannot leverage multiprocessors
- Entire process blocks when one thread blocks

# OS SUPPORT: APPROACH 2

## Kernel-level threads: One-to-one thread mapping

- OS provides each user-level thread with a kernel thread
- Each kernel thread scheduled independently
- Thread operations (creation, scheduling, synchronization) performed by OS

## Advantages

- Each kernel-level thread can run in parallel on a multiprocessor
- When one thread blocks, other threads from process can be scheduled

## Disadvantages

- Higher overhead for thread operations
- OS must scale well with increasing number of threads

# THREADS DEMO

# THREAD SCHEDULE #1

balance = balance + 1; balance at 0x9cd4

Registers are virtualized by OS;  
Each thread thinks it has own

**State:**  
0x9cd4: 100  
%eax: ?  
%rip = 0x195

process  
control  
blocks:

Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195



TI      0x195    mov 0x9cd4, %eax  
         0x19a    add \$0x1, %eax  
         0x19d    mov %eax, 0x9cd4A

What is state after instruction 0x195 completes?



# THREAD SCHEDULE #1

**State:**

0x9cd4: 100

%eax: 100

%rip = 0x19a

process  
control  
blocks:

Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195

TI →

0x195	mov 0x9cd4, %eax
0x19a	add \$0x1, %eax
0x19d	mov %eax, 0x9cd4A

What is state after instruction 0x19a completes?

# THREAD SCHEDULE #1

## State:

0x9cd4: 100

%eax: 101

%rip = 0x19d

process  
control  
blocks:

Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195

TI →

0x195 mov 0x9cd4, %eax

0x19a add \$0x1, %eax

0x19d mov %eax, 0x9cd4A

What is state after instruction 0x19d completes?

# THREAD SCHEDULE #1

## State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process  
control  
blocks:

Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195

TI →

```
0x195  mov 0x9cd4, %eax
0x19a  add $0x1, %eax
0x19d  mov %eax, 0x9cd4A
```

## Thread Context Switch

New contents of PCB and %eax and %rip?

# THREAD SCHEDULE #1

## State:

0x9cd4: 101

%eax: ?

%rip = 0x195

process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x1a2

Thread 2

%eax: ?  
%rip: 0x195

T2 →

0x195 mov 0x9cd4, %eax

0x19a add \$0x1, %eax

0x19d mov %eax, 0x9cd4A

What is state after instruction 0x195 completes?

# THREAD SCHEDULE #1

## State:

0x9cd4: 101  
%eax: 101  
%rip = 0x19a

process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x1a2

Thread 2

%eax: ?  
%rip: 0x195

T2 →  
0x195 mov 0x9cd4, %eax  
0x19a add \$0x1, %eax  
0x19d mov %eax, 0x9cd4A

What is state after instruction 0x19a completes?

# THREAD SCHEDULE #1

## State:

0x9cd4: 101  
%eax: 102  
%rip = 0x19d

process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x1a2

Thread 2

%eax: ?  
%rip: 0x195

T2 →

0x195	mov 0x9cd4, %eax
0x19a	add \$0x1, %eax
0x19d	mov %eax, 0x9cd4A

What is state after instruction 0x19d completes?

# THREAD SCHEDULE #1

## State:

0x9cd4: 102  
%eax: 102  
%rip = 0x1a2

process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x1a2

Thread 2

%eax: ?  
%rip: 0x195

0x195 mov 0x9cd4, %eax  
0x19a add \$0x1, %eax  
0x19d mov %eax, 0x9cd4A

T2 →

Desired Result!

**ANOTHER SCHEDULE**



## THREAD SCHEDULE #2

**State:**

0x9cd4: 100  
%eax: ?  
%rip = 0x195

process  
control  
blocks:

Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195



```
0x195  mov 0x9cd4, %eax
0x19a  add $0x1, %eax
0x19d  mov %eax, 0x9cd4A
```

## THREAD SCHEDULE #2

**State:**

0x9cd4: 100  
%eax: 100  
%rip = 0x19a

process  
control  
blocks:

Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195

TI →

0x195 mov 0x9cd4, %eax  
0x19a add \$0x1, %eax  
0x19d mov %eax, 0x9cd4A

## THREAD SCHEDULE #2

### State:

0x9cd4: 100  
%eax: 101  
%rip = 0x19d

process  
control  
blocks:

Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195



T1

0x195 mov 0x9cd4, %eax  
0x19a add \$0x1, %eax  
0x19d mov %eax, 0x9cd4A

Thread Context Switch before T1 executes 0x19d

## THREAD SCHEDULE #2

**State:**

0x9cd4: 100

%eax: ?

%rip = 0x195

process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x19d

Thread 2

%eax: ?  
%rip: 0x195

T2 →

0x195	mov 0x9cd4, %eax
0x19a	add \$0x1, %eax
0x19d	mov %eax, 0x9cd4A

## THREAD SCHEDULE #2

**State:**

0x9cd4: 100

%eax: 100

%rip = 0x19a

process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x19d

Thread 2

%eax: ?  
%rip: 0x195

T2 →

```
0x195  mov 0x9cd4, %eax
0x19a  add $0x1, %eax
0x19d  mov %eax, 0x9cd4A
```

## THREAD SCHEDULE #2

**State:**

0x9cd4: 100  
%eax: 101  
%rip = 0x19d

process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x19d

Thread 2

%eax: ?  
%rip: 0x195

T2 →

0x195 mov 0x9cd4, %eax  
0x19a add \$0x1, %eax  
0x19d mov %eax, 0x9cd4A

## THREAD SCHEDULE #2

**State:**

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x19d

Thread 2

%eax: ?  
%rip: 0x195

T2 →

```
0x195  mov 0x9cd4, %eax
0x19a  add $0x1, %eax
0x19d  mov %eax, 0x9cd4A
```

Thread Context Switch back to T1

## THREAD SCHEDULE #2

**State:**

0x9cd4: 101  
%eax: 101  
%rip = 0x19d

process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x19d

Thread 2

%eax: 101  
%rip: 0x1a2



0x195 mov 0x9cd4, %eax  
0x19a add \$0x1, %eax  
0x19d mov %eax, 0x9cd4A



## THREAD SCHEDULE #2

### State:

0x9cd4: 101  
%eax: 101  
%rip = 0x1a2

process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x1a2

Thread 2

%eax: 101  
%rip: 0x1a2

```
0x195  mov 0x9cd4, %eax
0x19a  add $0x1, %eax
0x19d  mov %eax, 0x9cd4A
```

TI →

## THREAD SCHEDULE #2

### State:

0x9cd4: 101  
%eax: 101  
%rip = 0x1a2

process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x1a2

Thread 2

%eax: 101  
%rip: 0x1a2

0x195 mov 0x9cd4, %eax  
0x19a add \$0x1, %eax  
0x19d mov %eax, 0x9cd4A

TI →



WRONG Result! Final value of balance is 101

# TIMELINE VIEW

Thread 1	Thread 2	Balance	T1-eax	T2-eax
mov 0x123, %eax		0	0	
			1	
add %0x1, %eax				
mov %eax, 0x123		1		
	mov 0x123, %eax			1
	add %0x2, %eax			3
	mov %eax, 0x123	3		

How much is added to shared variable?      3: correct!

# TIMELINE VIEW

Thread 1	Thread 2	Balance	T1-eax	T2-eax
mov 0x123, %eax		0	0	
add %0x1, %eax			1	
	mov 0x123, %eax			0
mov %eax, 0x123		1		
	add %0x2, %eax			2
	mov %eax, 0x123	2		

How much is added?

2: incorrect!

# TIMELINE VIEW

Thread 1	Thread 2	Balance	T1-eax	T2-eax
	mov 0x123, %eax	0		0
mov 0x123, %eax			0	
	add %0x2, %eax			2
add %0x1, %eax			1	
	mov %eax, 0x123	2		
mov %eax, 0x123		1		

How much is added?

1: incorrect!

# TIMELINE VIEW

Thread 1	Thread 2	Balance	T1-eax	T2-eax
	mov 0x123, %eax	0		0
	add %0x2, %eax			2
	mov %eax, 0x123	2		
mov 0x123, %eax			2	
add %0x1, %eax			3	
mov %eax, 0x123		3		

How much is added?

3: correct!

# TIMELINE VIEW

Thread 1	Thread 2	Balance	T1-eax	T2-eax
	mov 0x123, %eax	0		0
	add %0x2, %eax			2
mov 0x123, %eax			0	
add %0x1, %eax			1	
mov %eax, 0x123		1		
	mov %eax, 0x123	2		

How much is added?      2: incorrect!

# NON-DETERMINISM

Concurrency leads to non-deterministic results

- Different results even with same inputs
- race conditions

Whether bug manifests depends on CPU schedule!

How to program: imagine scheduler is malicious?!



# WHAT DO WE WANT?

Want 3 instructions to execute as an uninterruptable group

That is, we want them to be **atomic**

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

More general: **Need mutual exclusion** for **critical sections**

if thread A is in critical section C, thread B isn't  
(okay if other threads do unrelated work)

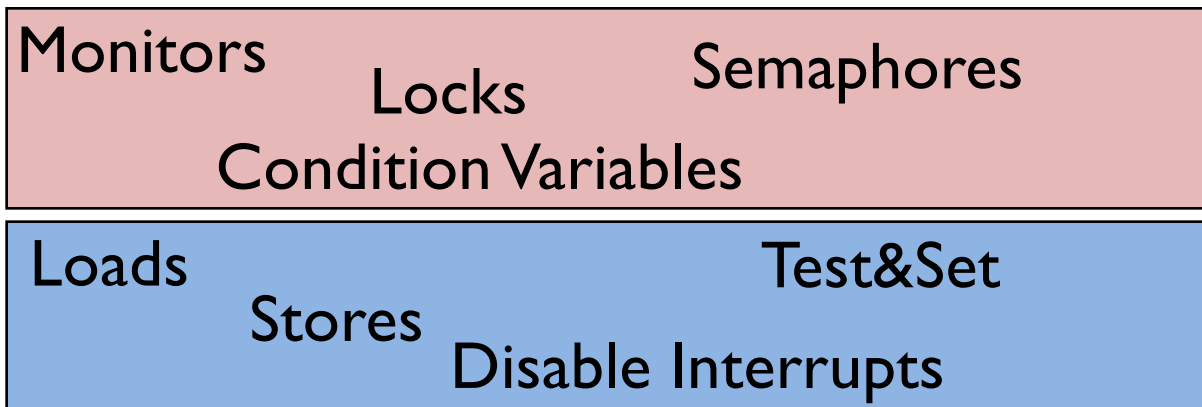
# SYNCHRONIZATION

Build higher-level synchronization primitives in OS

Operations that ensure correct ordering of instructions across threads

Use help from hardware

Motivation: Build them once and get them right



# CONCURRENCY SUMMARY

Concurrency is needed for high performance when using multiple cores

Threads are multiple execution streams within a single process or address space (share PID and address space, own registers and stack)

Context switches within a critical section can lead to non-deterministic bugs

# LOCKS

Goal: Provide mutual exclusion (**mutex**)

Allocate and Initialize

- **Pthread\_mutex\_t** mylock = PTHREAD\_MUTEX\_INITIALIZER;

Acquire

- Acquire exclusion access to lock;
- Wait if lock is not available (some other process in critical section)
- Spin or block (relinquish CPU) while waiting
- **Pthread\_mutex\_lock**(&mylock);

Release

- Release exclusive access to lock; let another process enter critical section
- **Pthread\_mutex\_unlock**(&mylock);

# THREADS DEMO2